# KISHORE B

AI / GenAI Engineer | Data Scientist | Python Developer | Data Engineer | Power BI Analyst
Chennai, India | blkishore.kk@gmail.com | +91-8610102348 | linkedin.com/in/kishore0408 | blkishore.github.io

## PROFESSIONAL SUMMARY

AI Engineer and Data Scientist with 3+ years of production experience building LLM-powered agentic systems, Retrieval-Augmented Generation (RAG) pipelines, machine learning models, enterprise ETL/ELT data pipelines, and BI dashboards for financial-services clients. Strong grounding in enterprise data platforms, distributed data processing, data quality, data governance, and data security, with CI/CD and DevOps practices applied across AI and data workloads. Proven, quantified outcomes: 40% reduction in AI retrieval latency, ~30% reduction in LLM hallucination rate, 60%+ manual-effort elimination through agentic automation, 25% improvement in AI pipeline throughput, and 50% faster BI report generation. Databricks Certified Data Engineer Associate, Microsoft Certified Fabric Data Engineer Associate, and Microsoft Certified Power BI Data Analyst Associate, with additional Pega CSSA/CSA architecture certifications. Works daily with Claude Code, GitHub Copilot, and Cursor as AI-assisted development partners.

## CORE TECHNICAL SKILLS

**Generative AI & LLM Engineering:** AI/ML, Large Language Models (LLMs), OpenAI APIs, Retrieval-Augmented Generation (RAG), Agentic AI Workflows, Prompt Engineering, System-Prompt & Guardrail Design, Function/Tool Calling, LLM Evaluation & Hallucination Reduction, Embedding Generation, Semantic Search, Context Management, LangChain, LlamaIndex, Graph RAG, Hugging Face, NLP

**Agentic Orchestration & Vector Databases:** LangChain and LlamaIndex (production); LangGraph, CrewAI, AutoGen (multi-agent orchestration); multi-step agentic workflow design, tool-use chaining, stateful agent context management; Weaviate, Pinecone, FAISS, Qdrant

**Machine Learning & Data Science:** Scikit-learn, PyTorch, TensorFlow, Model Fine-tuning (LoRA/QLoRA), Classification, Regression, Clustering, Anomaly Detection, Statistical Modelling, Hypothesis Testing, Feature Engineering, EDA, Predictive Analytics, Model Monitoring & Drift Tracking

**Data Engineering & Enterprise Data Platforms:** ETL/ELT Pipeline Design, Data Engineering, Data Modelling, Lakehouse Architecture, Distributed Data Processing (PySpark), Databricks (certified), Microsoft Fabric, Azure Data Factory (ADF), Azure Data Lake Storage (ADLS), Informatica PowerCenter, Data Quality & Validation, Data Governance, Data Security, Performance Optimization

**Cloud Platforms:** Microsoft Azure / Microsoft Fabric, Azure Cloud Deployments; AWS (Amazon S3, AWS Glue, AWS Lambda)

**Databases & SQL:** SQL (advanced, query optimisation), Stored Procedures, PostgreSQL, MySQL, NoSQL Databases, JSON

**Programming & Backend:** Python, Object-Oriented Programming (OOP), FastAPI, Django, REST API Design & Development, Java, JavaScript, React.js, C#/.NET (working exposure)

**DevOps, MLOps & LLMOps:** CI/CD (GitHub Actions, Jenkins), DevOps, Container Management, MLflow, LangSmith, Prometheus/Grafana, Git, Agile/Scrum

**BI & Visualization:** Power BI (DAX, Data Modelling, KPI Design, Forecasting, Drill-through Reporting), Tableau, Streamlit, Dashboard & Monitoring UI Development

**Automation & Workflow Orchestration:** n8n (visual workflow orchestration), End-to-End AI Pipeline Automation, Event-driven Automation

**AI-Assisted Development:** Claude Code, GitHub Copilot, Cursor — used daily for pair-programming, code review, and rapid prototyping of AI features

## PROFESSIONAL EXPERIENCE

**AI / Python Developer — AI-Focused Role (AIG Project)** **2024 – Present**
*Tata Consultancy Services (TCS)*

- **RAG Pipeline Architecture:** Architected end-to-end Retrieval-Augmented Generation pipelines using OpenAI APIs and LangChain, reducing information-retrieval latency by 40% and improving answer accuracy across enterprise knowledge bases spanning 100K+ records.
- **Agentic Workflow Design:** Designed multi-step agentic AI workflows incorporating tool use, LLM orchestration, and persistent context management — automating complex enterprise tasks end-to-end in a production environment.
- **Prompt Engineering & Guardrails:** Wrote and iteratively refined system prompts, instructions, and safety guardrails, cutting hallucination rates by approximately 30% across production LLM workflows.
- **Backend & API Development:** Built FastAPI backend services and REST APIs integrating LLM systems with enterprise data sources; optimised AI pipeline throughput by 25% through systematic performance optimization and tuning.
- **Embeddings & Semantic Search:** Engineered embedding generation and semantic search pipelines to power accurate document retrieval for downstream LLM consumption.
- **Observability:** Designed Streamlit monitoring dashboards for real-time AI pipeline observability, LLM output evaluation, and stakeholder interaction.

**Software Engineer — Enterprise Applications (Corebridge Financial)** **2024 – Present**
*Tata Consultancy Services (TCS)*

- **Enterprise Application Development:** Developed enterprise-grade, object-oriented Java applications supporting 500K+ daily financial transactions, ensuring reliability and performance at scale.

- **Data Transformation & Reporting:** Applied advanced SQL (including stored procedures) for data transformation, validation, and reporting, maintaining 99.8% data reliability across critical financial workflows.
- **Production Support:** Resolved 50+ critical production incidents through structured root-cause analysis, minimising downtime and business impact.
- **Agile Delivery:** Collaborated within Agile/Scrum teams using Git for version control, delivering 15+ feature releases on schedule.

**Python & AI Developer (Freelance)** **2023 – 2024**

*Self-Employed*

- **AI Automation Pipelines:** Designed AI automation pipelines using n8n, LLMs, and OpenAI APIs, eliminating 60%+ of manual operations for business clients through agentic, multi-step workflows.
- **API Integration:** Integrated 10+ external REST APIs into NLP-based processing pipelines to extend automation coverage across client systems.
- **AI Copilot Interfaces:** Built Streamlit-based AI copilot interfaces enabling non-technical users to interact directly with AI-generated outputs.

## KEY PROJECTS

**AI Automation & Agentic System — n8n + OpenAI + RAG + LangChain**

- Built a production-grade agentic AI system combining RAG pipelines with embedding-based vector retrieval (85%+ retrieval accuracy) and LLM-powered classification and summarisation, reducing manual effort by 60%.
- Developed Streamlit monitoring dashboards for real-time LLM output evaluation, pipeline observability, and stakeholder interaction.

**Power BI Analytics Dashboard — Financial Domain**

- Built DAX-driven KPI dashboards with trend analysis, drill-through reporting, and AI-assisted forecasting.
- Connected multiple data sources via SQL-based ETL models, reducing report-generation time by 50%.

**ML Predictive Analytics & Anomaly Detection System**

- Delivered an end-to-end ML pipeline — EDA, feature engineering, classification/anomaly detection modelling, and Azure deployment.
- Exposed model predictions via a FastAPI inference endpoint with post-deployment model monitoring.

## CERTIFICATIONS

- Databricks Certified Data Engineer Associate
- Microsoft Certified: Fabric Data Engineer Associate
- Microsoft Certified: Power BI Data Analyst Associate
- Pega Certified Senior System Architect (CSSA)
- Pega Certified System Architect (CSA)

## EDUCATION

**Bachelor of Engineering (B.E.) — Electronics and Communication Engineering**

K. Ramakrishnan College of Engineering, Chennai